

COMPSCIX 415.2 Homework 7

Vishnu Vardhan

3/17/2018

Contents

Exercise 1	1
Answer	1
Exercise 2	3
Answer	4
Exercise 3	4
Answer	4
Exercise 4	5
Answer	5
Exercise 5	6
Answer	6
Exercise 6 (OPTIONAL - won't be graded)	6
Answer	7
Exercise 7	7
Answer	7

Data sources from Kaggle: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/leaderboard>

Data has not been uploaded into git

Exercise 1

Load the train.csv dataset into R. How many observations and columns are there?

Answer

There are 81 variables, and 1460 observations

```
train <- read_csv('train.csv')

## Parsed with column specification:
## cols(
##   .default = col_character(),
##   Id = col_integer(),
##   MSSubClass = col_integer(),
##   LotFrontage = col_integer(),
##   LotArea = col_integer(),
##   OverallQual = col_integer(),
##   OverallCond = col_integer(),
##   YearBuilt = col_integer(),
##   YearRemodAdd = col_integer(),
##   MasVnrArea = col_integer(),
##   BsmtFinSF1 = col_integer(),
##   BsmtFinSF2 = col_integer(),
##   BsmtUnfSF = col_integer(),
```

```
## TotalBsmtSF = col_integer(),
## `1stFlrSF` = col_integer(),
## `2ndFlrSF` = col_integer(),
## LowQualFinSF = col_integer(),
## GrLivArea = col_integer(),
## BsmtFullBath = col_integer(),
## BsmtHalfBath = col_integer(),
## FullBath = col_integer()
## # ... with 18 more columns
## )

## See spec(...) for full column specifications.
```

```
glimpse(train)
```

```
## Observations: 1,460
## Variables: 81
## $ Id <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 1...
## $ MSSubClass <int> 60, 20, 60, 70, 60, 50, 20, 60, 50, 190, 20, 60,...
## $ MSZoning <chr> "RL", "RL", "RL", "RL", "RL", "RL", "RL", "RL", ...
## $ LotFrontage <int> 65, 80, 68, 60, 84, 85, 75, NA, 51, 50, 70, 85, ...
## $ LotArea <int> 8450, 9600, 11250, 9550, 14260, 14115, 10084, 10...
## $ Street <chr> "Pave", "Pave", "Pave", "Pave", "Pave", "Pave", "Pave", ...
## $ Alley <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ LotShape <chr> "Reg", "Reg", "IR1", "IR1", "IR1", "IR1", "Reg",...
## $ LandContour <chr> "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl", "Lvl",...
## $ Utilities <chr> "AllPub", "AllPub", "AllPub", "AllPub", "AllPub"...
## $ LotConfig <chr> "Inside", "FR2", "Inside", "Corner", "FR2", "Ins...
## $ LandSlope <chr> "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl", "Gtl",...
## $ Neighborhood <chr> "CollgCr", "Veenker", "CollgCr", "Crawfor", "NoR...
## $ Condition1 <chr> "Norm", "Feedr", "Norm", "Norm", "Norm", "Norm",...
## $ Condition2 <chr> "Norm", "Norm", "Norm", "Norm", "Norm", "Norm", ...
## $ BldgType <chr> "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", "1Fam", ...
## $ HouseStyle <chr> "2Story", "1Story", "2Story", "2Story", "2Story"...
## $ OverallQual <int> 7, 6, 7, 7, 8, 5, 8, 7, 7, 5, 5, 9, 5, 7, 6, 7, ...
## $ OverallCond <int> 5, 8, 5, 5, 5, 5, 5, 6, 5, 6, 5, 5, 6, 5, 5, 8, ...
## $ YearBuilt <int> 2003, 1976, 2001, 1915, 2000, 1993, 2004, 1973, ...
## $ YearRemodAdd <int> 2003, 1976, 2002, 1970, 2000, 1995, 2005, 1973, ...
## $ RoofStyle <chr> "Gable", "Gable", "Gable", "Gable", "Gable", "Ga...
## $ RoofMatl <chr> "CompShg", "CompShg", "CompShg", "CompShg", "Com...
## $ Exterior1st <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Sdng", "Vin...
## $ Exterior2nd <chr> "VinylSd", "MetalSd", "VinylSd", "Wd Shng", "Vin...
## $ MasVnrType <chr> "BrkFace", "None", "BrkFace", "None", "BrkFace",...
## $ MasVnrArea <int> 196, 0, 162, 0, 350, 0, 186, 240, 0, 0, 0, 286, ...
## $ ExterQual <chr> "Gd", "TA", "Gd", "TA", "Gd", "TA", "Gd", "TA", ...
## $ ExterCond <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ Foundation <chr> "PConc", "CBlock", "PConc", "BrkTil", "PConc", "...
## $ BsmtQual <chr> "Gd", "Gd", "Gd", "TA", "Gd", "Gd", "Ex", "Gd", ...
## $ BsmtCond <chr> "TA", "TA", "TA", "Gd", "TA", "TA", "TA", "TA", ...
## $ BsmtExposure <chr> "No", "Gd", "Mn", "No", "Av", "No", "Av", "Mn", ...
## $ BsmtFinType1 <chr> "GLQ", "ALQ", "GLQ", "ALQ", "GLQ", "GLQ", "GLQ",...
## $ BsmtFinSF1 <int> 706, 978, 486, 216, 655, 732, 1369, 859, 0, 851,...
## $ BsmtFinType2 <chr> "Unf", "Unf", "Unf", "Unf", "Unf", "Unf", "Unf",...
## $ BsmtFinSF2 <int> 0, 0, 0, 0, 0, 0, 0, 32, 0, 0, 0, 0, 0, 0, 0, ...
## $ BsmtUnfSF <int> 150, 284, 434, 540, 490, 64, 317, 216, 952, 140,...
```

```

## $ TotalBsmtSF <int> 856, 1262, 920, 756, 1145, 796, 1686, 1107, 952,...
## $ Heating <chr> "GasA", "GasA", "GasA", "GasA", "GasA", "GasA", ...
## $ HeatingQC <chr> "Ex", "Ex", "Ex", "Gd", "Ex", "Ex", "Ex", "Ex", ...
## $ CentralAir <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ Electrical <chr> "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SBrkr", "SB...
## $ `1stFlrSF` <int> 856, 1262, 920, 961, 1145, 796, 1694, 1107, 1022...
## $ `2ndFlrSF` <int> 854, 0, 866, 756, 1053, 566, 0, 983, 752, 0, 0, ...
## $ LowQualFinSF <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ GrLivArea <int> 1710, 1262, 1786, 1717, 2198, 1362, 1694, 2090, ...
## $ BsmtFullBath <int> 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, ...
## $ BsmtHalfBath <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ FullBath <int> 2, 2, 2, 1, 2, 1, 2, 2, 2, 1, 1, 3, 1, 2, 1, 1, ...
## $ HalfBath <int> 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, ...
## $ BedroomAbvGr <int> 3, 3, 3, 3, 4, 1, 3, 3, 2, 2, 3, 4, 2, 3, 2, 2, ...
## $ KitchenAbvGr <int> 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, 1, ...
## $ KitchenQual <chr> "Gd", "TA", "Gd", "Gd", "Gd", "TA", "Gd", "TA", ...
## $ TotRmsAbvGrd <int> 8, 6, 6, 7, 9, 5, 7, 7, 8, 5, 5, 11, 4, 7, 5, 5,...
## $ Functional <chr> "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", "Typ", ...
## $ Fireplaces <int> 0, 1, 1, 1, 1, 0, 1, 2, 2, 2, 0, 2, 0, 1, 1, 0, ...
## $ FireplaceQu <chr> NA, "TA", "TA", "Gd", "TA", NA, "Gd", "TA", "TA"...
## $ GarageType <chr> "Attchd", "Attchd", "Attchd", "Detchd", "Attchd"...
## $ GarageYrBlt <int> 2003, 1976, 2001, 1998, 2000, 1993, 2004, 1973, ...
## $ GarageFinish <chr> "RFn", "RFn", "RFn", "Unf", "RFn", "Unf", "RFn",...
## $ GarageCars <int> 2, 2, 2, 3, 3, 2, 2, 2, 2, 1, 1, 3, 1, 3, 1, 2, ...
## $ GarageArea <int> 548, 460, 608, 642, 836, 480, 636, 484, 468, 205...
## $ GarageQual <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ GarageCond <chr> "TA", "TA", "TA", "TA", "TA", "TA", "TA", "TA", ...
## $ PavedDrive <chr> "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y", "Y"...
## $ WoodDeckSF <int> 0, 298, 0, 0, 192, 40, 255, 235, 90, 0, 0, 147, ...
## $ OpenPorchSF <int> 61, 0, 42, 35, 84, 30, 57, 204, 0, 4, 0, 21, 0, ...
## $ EnclosedPorch <int> 0, 0, 0, 272, 0, 0, 0, 228, 205, 0, 0, 0, 0, 0, ...
## $ `3SsnPorch` <int> 0, 0, 0, 0, 0, 320, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ ScreenPorch <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 176, 0, 0, 0...
## $ PoolArea <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ PoolQC <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ Fence <chr> NA, NA, NA, NA, NA, "MnPrv", NA, NA, NA, NA, NA,...
## $ MiscFeature <chr> NA, NA, NA, NA, NA, "Shed", NA, "Shed", NA, NA, ...
## $ MiscVal <int> 0, 0, 0, 0, 0, 700, 0, 350, 0, 0, 0, 0, 0, 0, 0,...
## $ MoSold <int> 2, 5, 9, 2, 12, 10, 8, 11, 4, 1, 2, 7, 9, 8, 5, ...
## $ YrSold <int> 2008, 2007, 2008, 2006, 2008, 2009, 2007, 2009, ...
## $ SaleType <chr> "WD", "WD", "WD", "WD", "WD", "WD", "WD", "WD", ...
## $ SaleCondition <chr> "Normal", "Normal", "Normal", "Abnorml", "Normal...
## $ SalePrice <int> 208500, 181500, 223500, 140000, 250000, 143000, ...

```

Exercise 2

Normally at this point you would spend a few days on EDA, but for this homework we will get right to fitting some linear regression models.

Our first step is to randomly split the data into train and test datasets. We will use a 70/30 split. There is an R package that will do the split for you, but let's get some more practice with R and do it ourselves by filling in the blanks in the code below.

Answer

```
# load packages
library(tidyverse)
library(broom)
# When taking a random sample, it is often useful to set a seed so that
# your work is reproducible. Setting a seed will guarantee that the same
# random sample will be generated every time, so long as you always set the
# same seed beforehand
set.seed(29283)
# This data already has an Id column which we can make use of.
# Let's create our training set using sample_frac. Fill in the blank.
train_set <- train %>% sample_frac(size=0.7, replace = FALSE)
# let's create our testing set using the Id column. Fill in the blanks.
test_set <- train %>% filter(!(Id %in% train_set$Id))
```

Exercise 3

Our target is called SalePrice. First, we can fit a simple regression model consisting of only the intercept (the average of SalePrice). Fit the model and then use the broom package to

- take a look at the coefficient,
- compare the coefficient to the average value of SalePrice, and
- take a look at the R-squared.

Use the code below and fill in the blanks.

Answer

The coefficient of to the average sale price is almost the same as the mean sales price.

With an R-Squared value = 0, none of the variation in the data is explained by the mean.

```
# Fit a model with intercept only
mod_0 <- lm(SalePrice ~ 1, data = train_set)
# Double-check that the average SalePrice is equal to our model's coefficient
sprintf ("Mean = %f", mean(train_set$SalePrice))
```

```
## [1] "Mean = 182175.954990"
```

```
tidy(mod_0)
```

```
##           term estimate std.error statistic p.value
## 1 (Intercept)  182176   2492.072    73.10222      0
```

```
# Check the R-squared
glance(mod_0)
```

```
##   r.squared adj.r.squared   sigma statistic p.value df   logLik   AIC
## 1         0             0 79668.37      NA      NA  1 -12983.57 25971.13
##           BIC      deviance df.residual
## 1 25980.99 6.480338e+12          1021
```

Exercise 4

Now fit a linear regression model using GrLivArea, OverallQual, and Neighborhood as the features. Don't forget to look at data_description.txt to understand what these variables mean. Ask yourself these questions before fitting the model:

- What kind of relationship will these features have with our target?
- Can the relationship be estimated linearly?
- Are these good features, given the problem we are trying to solve?

After fitting the model, output the coefficients and the R-squared using the broom package.

Answer these questions:

- How would you interpret the coefficients on GrLivArea and OverallQual?
- How would you interpret the coefficient on NeighborhoodBrkSide?
- Are the features significant?
- Are the features practically significant?
- Is the model a good fit (to the training set)?

Answer

```
mod_1 <- lm(SalePrice ~ GrLivArea + OverallQual + Neighborhood, data = train_set)
tidy(mod_1)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	-45017.87483	12933.341808	-3.4807612	5.216927e-04
## 2	GrLivArea	62.77735	3.006033	20.8837885	1.337222e-80
## 3	OverallQual	21692.23178	1353.714104	16.0242342	1.389020e-51
## 4	NeighborhoodBlueste	-38288.88063	36531.907177	-1.0480942	2.948497e-01
## 5	NeighborhoodBrDale	-43314.05372	14524.693991	-2.9820975	2.932566e-03
## 6	NeighborhoodBrkSide	-14064.37052	11318.850018	-1.2425618	2.143221e-01
## 7	NeighborhoodClearCr	27839.00662	13561.346871	2.0528202	4.035110e-02
## 8	NeighborhoodCollgCr	4297.67432	10372.304467	0.4143413	6.787135e-01
## 9	NeighborhoodCrawfor	7423.05573	11371.511784	0.6527765	5.140512e-01
## 10	NeighborhoodEdwards	-15284.11495	10994.287187	-1.3901870	1.647830e-01
## 11	NeighborhoodGilbert	-8357.55930	10894.173472	-0.7671586	4.431692e-01
## 12	NeighborhoodIDOTRR	-32689.43085	12603.712743	-2.5936350	9.636216e-03
## 13	NeighborhoodMeadowV	-14446.06504	14190.148622	-1.0180348	3.089089e-01
## 14	NeighborhoodMitchel	1922.31487	11788.608170	0.1630655	8.705000e-01
## 15	NeighborhoodNAMES	-7719.67883	10375.956174	-0.7439969	4.570540e-01
## 16	NeighborhoodNoRidge	47685.16790	12567.432633	3.7943444	1.569690e-04
## 17	NeighborhoodNPkVill	-20240.71145	16548.664867	-1.2231024	2.215806e-01
## 18	NeighborhoodNridgHt	63872.80848	10880.456671	5.8704161	5.917964e-09
## 19	NeighborhoodNWAmes	-12279.33299	11047.502893	-1.1115030	2.666204e-01
## 20	NeighborhoodOldTown	-36107.07577	10849.170903	-3.3280954	9.064637e-04
## 21	NeighborhoodSawyer	-4121.92502	11252.369778	-0.3663162	7.142070e-01
## 22	NeighborhoodSawyerW	-5391.97074	11230.758221	-0.4801075	6.312565e-01
## 23	NeighborhoodSomerst	18700.96725	10772.212794	1.7360377	8.286672e-02
## 24	NeighborhoodStoneBr	65712.45881	12745.312907	5.1558137	3.045915e-07
## 25	NeighborhoodSWISU	-45451.86707	13564.792586	-3.3507233	8.363074e-04
## 26	NeighborhoodTimber	27925.08619	11985.325857	2.3299397	2.000859e-02
## 27	NeighborhoodVeenker	54913.12768	16521.075497	3.3238228	9.203087e-04

- How would you interpret the coefficients on GrLivArea and OverallQual?

GrLivArea coefficient of 62.77 says that all things being equal, the per-sq-ft charge is about 62 units. The OverallQual parameter specifies that the quality of the house increases the value by 21,692 at each level of quality.

- How would you interpret the coefficient on NeighborhoodBrkSide?

The co-efficient implies that being in that neighborhood penalizes a home (by -14,064), and its value drops. But because the p-value is 0.214, it is not a significant predictor.

- Are the features significant?

Not all p-values are significant. Since p-values have to be less than 0.05, many of the neighborhood p-values are not significant.

- Are the features practically significant?

Yes, the features are practically significant, since we know that these factors play critically into a home's value. There are other factors that could improve our prediction.

```
glance(mod_1)
```

```
##   r.squared adj.r.squared   sigma statistic p.value df   logLik   AIC
## 1 0.8099927   0.8050277 35178.1 163.1401      0 27 -12134.95 24325.91
##           BIC      deviance df.residual
## 1 24463.93 1.231311e+12          995
```

- Is the model a good fit (to the training set)?

Yes, the model is a good-fit, because the adjusted R-squared is about 0.8. I.e the model explains 80% of the variance, which is much better than a random guess. It could be better

Exercise 5

Evaluate the model on test_set using the root mean squared error (RMSE). Use the predict function to get the model predictions for the testing set. Recall that RMSE is the square root of the mean of the squared errors:

Hint: use the sqrt() and mean() functions:

```
test_predictions <- predict(NAME_OF_YOUR_MODEL_HERE, newdata = test_set) rmse <-
sqrt(mean((____ - ____)^2))
```

Answer

```
test_predictions <- predict(mod_1, newdata = test_set)
rmse <- sqrt(mean((test_predictions - test_set$SalePrice)^2))
rmse
```

```
## [1] 41915.27
```

The root mean square value is 41,915. Given the summary stats below for the sale price, the value looks reasonable

```
summary(train$SalePrice)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.   Max.
##  34900 129975  163000  180921  214000  755000
```

Exercise 6 (OPTIONAL - won't be graded)

Feel free to play around with linear regression. Add some other features and see how the model results change. Test the model on test_set to compare the RMSE's.

Answer

Seems like i was able to improve from the 41K RMS value to about 36K, by including 1. Year built 2. Interactions between BldgType and Neighborhood 3. Interactions between Neighborhood and OverallQuality

```
mod_4 <- lm(SalePrice ~ GrLivArea +
            OverallQual +
            Neighborhood +
            YearBuilt +
            BldgType:Neighborhood +
            Neighborhood:OverallQual,
            data = train_set)
test_predictions <- predict(mod_4, newdata = test_set)

## Warning in predict.lm(mod_4, newdata = test_set): prediction from a rank-
## deficient fit may be misleading

rmse <- sqrt(mean((test_predictions - test_set$SalePrice)^2))
rmse

## [1] 36646.79
```

Exercise 7

One downside of the linear model is that it is sensitive to unusual values because the distance incorporates a squared term. Fit a linear model to the simulated data below, and visualise the results. Rerun a few times to generate different simulated datasets.

What do you notice about the model?

```
sim1a <- tibble(
  x = rep(1:10, each = 3),
  y = x * 1.5 + 6 + rt(length(x), df = 2)
)
```

Answer

As can be seen in the code below, the r-squared values can vary pretty dramatically, with most of the values being above 0.75

```
df <- tibble(run = integer(), rsquared = double())
for (i in c(1:1000)) {
  sim1a <- tibble(
    x = rep(1:10, each = 3),
    y = x * 1.5 + 6 + rt(length(x), df = 2)
  )
  mod_2 <- lm(y ~ x, data=sim1a)
  results <- glance(mod_2)$adj.r.squared
  df <- add_row(df, run = i, rsquared = results)
}

ggplot(data = df, mapping=aes(x=rsquared)) + geom_histogram()

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

